

Q1

$$a) P(x|c, \theta) = \prod_{i=1}^N P(x_j^{(i)} | c, \theta_{jc}) = \prod_{i=1}^N \theta_{jc}^{x_j^{(i)}} (1 - \theta_{jc})^{(1 - x_j^{(i)})}$$

$$\begin{aligned} \log P(x|c, \theta) &= \sum_{i=1}^N \log(\theta_{jc}^{x_j^{(i)}} (1 - \theta_{jc})^{(1 - x_j^{(i)})}) \\ &= \sum_{i=1}^N (x_j^{(i)} \log \theta_{jc} + (1 - x_j^{(i)}) \log (1 - \theta_{jc})) \end{aligned}$$

$$\begin{aligned} \frac{\partial}{\partial \theta_{jc}} \log P(x|c, \theta) &= \sum_{i=1}^N \left(\frac{x_j^{(i)}}{\theta_{jc}} + \frac{(1 - x_j^{(i)})}{(-1)(1 - \theta_{jc})} \right) = \sum_{i=1}^N \frac{x_j^{(i)}}{\theta_{jc}} + \sum_{i=1}^N \frac{1 - x_j^{(i)}}{\theta_{jc} - 1} \\ &= \frac{1}{\theta_{jc}} \sum_{i=1}^N x_j^{(i)} + \frac{1}{\theta_{jc} - 1} \sum_{i=1}^N (1 - x_j^{(i)}) \\ &= \frac{1}{\theta_{jc}} N_{j, x_j=1} + \frac{1}{\theta_{jc} - 1} N_{j, x_j=0} \end{aligned}$$

where $N_{j, x_j=1}$ is the total number of images with pixel x_j classified as 1.

$N_{j, x_j=0}$ is the total number of images with pixel x_j classified as 0.

Setting the derivative to zero:

$$\begin{aligned} 0 &= \frac{1}{\theta_{jc}} N_{j, x_j=1} + \frac{1}{\theta_{jc} - 1} N_{j, x_j=0} \\ &= \frac{\hat{\theta}_{jc} - 1}{\hat{\theta}_{jc}(\hat{\theta}_{jc} - 1)} N_{j, x_j=1} + \frac{\hat{\theta}_{jc}}{\hat{\theta}_{jc}(\hat{\theta}_{jc} - 1)} N_{j, x_j=0} \\ &= \hat{\theta}_{jc} (N_{j, x_j=1} + N_{j, x_j=0}) - N_{j, x_j=1} \\ N_{j, x_j=1} &= \hat{\theta}_{jc} \cdot N_j \\ \hat{\theta}_{jc} &= \frac{N_{j, x_j=1}}{N_j} \end{aligned}$$

$$P(t^{(i)}|\pi) = \prod_{j=0}^q \pi_j t_j^{(i)}$$

$$\log P(t^{(i)}|\pi) = \sum_{j=0}^q \log(\pi_j t_j^{(i)})$$

$$= \sum_{j=0}^q t_j^{(i)} \log \pi_j$$

$$= \sum_{j=0}^8 t_j^{(i)} \log \pi_j +$$

$$t_9^{(i)} \log \left(1 - \sum_{j=0}^8 \pi_j\right)$$

$$\frac{\partial}{\partial \pi_j} \log P(t^{(i)}|\pi) = \sum_{j=0}^8 \frac{t_j^{(i)}}{\pi_j} - t_9^{(i)} \frac{1}{1 - \sum_{j=0}^8 \pi_j}$$

$$= \frac{\sum_{j=0}^8 t_j^{(i)}}{\sum_{j=0}^8 \pi_j} - \frac{t_9^{(i)}}{1 - \sum_{j=0}^8 \pi_j}$$

Setting derivative to zero:

$$\frac{t_9^{(i)}}{1 - \sum_{j=0}^8 \pi_j} = \frac{\sum_{j=0}^8 t_j^{(i)}}{\sum_{j=0}^8 \pi_j}$$

$$t_9^{(i)} \cdot \sum_{j=0}^8 \pi_j = \sum_{j=0}^8 t_j^{(i)} \cdot \left(1 - \sum_{j=0}^8 \pi_j\right)$$

$$\frac{\sum_{j=0}^8 \pi_j}{1 - \sum_{j=0}^8 \pi_j} = \frac{\sum_{j=0}^8 t_j^{(u)}}{t_q^{(u)}}$$

$$t_q^{(u)} \cdot \sum_{j=0}^8 \pi_j = \sum_{j=0}^8 t_j^{(u)} - \left(\sum_{j=0}^8 t_j^{(u)} \cdot \sum_{j=0}^8 \pi_j \right)$$

$$\sum_{j=0}^9 t_j^{(u)} \cdot \sum_{j=0}^8 \pi_j = \sum_{j=0}^8 t_j^{(u)}$$

$$\sum_{j=0}^8 \pi_j = \frac{\sum_{j=0}^8 t_j^{(u)}}{\sum_{j=0}^9 t_j^{(u)}}$$

$$\pi_q = 1 - \sum_{j=0}^8 \pi_j = \frac{t_q^{(u)}}{\sum_{j=0}^9 t_j^{(u)}}$$

$$\pi_j^* = \frac{\sum_{i=1}^N t_j^{(u)}}{\sum_{i=1}^N \sum_{j=0}^9 t_j^{(u)}}$$

$$b) p(t|x, \theta, \pi) \propto \prod_{i=1}^N p(t^{(i)}|\pi) p(x^{(i)}|t^{(i)}, \theta)$$

$$\mathcal{L}(\theta) = \prod_{i=1}^N p(t^{(i)}|\pi) \prod_{j=1}^{784} (\theta_{jt^{(i)}})^{x_j^{(i)}} (1 - \theta_{jt^{(i)}})^{(1-x_j^{(i)})}$$

$$\log \mathcal{L}(\theta) = \sum_{i=1}^N \left(\log \pi_{t^{(i)}} + \sum_{j=1}^{784} \left(x_j^{(i)} \log \theta_{jt^{(i)}} + (1 - x_j^{(i)}) \log (1 - \theta_{jt^{(i)}}) \right) \right)$$

for a single image:

$$\log \mathcal{L}(\theta) = \log \pi_{t^{(i)}} + \sum_{j=1}^{784} \left(x_j^{(i)} \log \theta_{jt^{(i)}} + (1 - x_j^{(i)}) \log (1 - \theta_{jt^{(i)}}) \right)$$


c) ./naive_bayes.py:173: RuntimeWarning: divide by zero encountered in log

```
sum_img_pixels = np.matmul(images, np.log(theta))
./naive_bayes.py:173: RuntimeWarning: invalid value encountered in
matmul
```

```
sum_img_pixels = np.matmul(images, np.log(theta))
Average log-likelihood for MLE is nan
```

According to the RuntimeWarning, division by zero will happen when calculating the average log-likelihood for MLE, which results in nan.

d) Plotting MLE estimators for 10 classes



$$e) \hat{\theta}_{MAP} = \underset{\theta}{\operatorname{argmax}} \log p(\theta) + \log p(D|\theta) \quad \theta \sim \text{Beta}(3, 3)$$

$$\begin{aligned} L(\theta) &= \log p(\theta_{jc}) + \sum_{i=1}^N \log(p(x_j^{(i)} | c, \theta_{jc})) \\ &= \log(\theta_{jc}^2 (1-\theta_{jc})^2) + \sum_{i=1}^N \log(\theta_{jc}^{x_j^{(i)}} (1-\theta_{jc})^{(1-x_j^{(i)})}) \\ &= 2\log\theta_{jc} + 2\log(1-\theta_{jc}) + \sum_{i=1}^N x_j^{(i)} \log\theta_{jc} + \sum_{i=1}^N (1-x_j^{(i)}) \log(1-\theta_{jc}) \\ &= 2\log\theta_{jc} + 2\log(1-\theta_{jc}) + N_{j, x_j=1} \log\theta_{jc} + N_{j, x_j=0} \log(1-\theta_{jc}) \\ &= \log\theta_{jc} \cdot (2 + N_{j, x_j=1}) + \log(1-\theta_{jc}) \cdot (2 + N_{j, x_j=0}) \end{aligned}$$

$$\frac{\partial L(\theta)}{\partial \theta_{jc}} = \frac{2 + N_{j, x_j=1}}{\theta_{jc}} + \frac{2 + N_{j, x_j=0}}{\theta_{jc} - 1}$$

setting the derivative to zero:

$$\frac{2 + N_{j, x_j=1}}{\theta_{jc}} = \frac{2 + N_{j, x_j=0}}{1 - \theta_{jc}}$$

$$\begin{aligned} 2 - 2\theta_{jc} + N_{j, x_j=1} - \theta_{jc} \cdot N_{j, x_j=1} \\ \parallel \\ 2\theta_{jc} + \theta_{jc} \cdot N_{j, x_j=0} \end{aligned}$$

$$2 + N_{j, x_j=1} = 4\theta_{jc} + \theta_{jc} \cdot (N_{j, x_j=1} + N_{j, x_j=0})$$

$$2 + N_{j, x_j=1} = \theta_{jc}(4 + N_j)$$

$$\hat{\theta}_{jc} = \frac{N_{j, x_j=1} + 2}{N_j + 4}$$


f)

Average log-likelihood for MAP is -100.39600482438387

Training accuracy for MAP is 0.6915833333333333

Test accuracy for MAP is 0.684

g) Plotting MAP estimators for 10 classes



Code for Q1:

```
def train_mle_estimator(train_images, train_labels):
    """ Inputs: train_images, train_labels
        Returns the MLE estimators theta_mle and pi_mle """

    # YOU NEED TO WRITE THIS PART
    #train_images.shape = (60000,784)
    #train_labels.shape = (60000,10)
    num_classes = 10
    num_img = train_images.shape[0]
    pixel_per_img = train_images.shape[1]
    train_data = np.where(train_images>0.5,1.0,0.0)

    num_pixel_in_class = np.matmul(train_labels.T,train_data)
    img_in_class = np.sum(train_labels.T,axis=1)
    img_in_class= img_in_class[:,np.newaxis]
    for i in range(0, num_classes):
        num_pixel_in_class[i] = num_pixel_in_class[i]/img_in_class[i][0]
    theta_mle = num_pixel_in_class.T
    # print("theta_mle",theta_mle.shape)

    img_in_class_vec = np.sum(train_labels,axis=0)
    sum_target_all_img = np.sum(np.sum(train_labels,axis=1),axis=0)
    pi_mle = img_in_class_vec/sum_target_all_img

    return theta_mle, pi_mle
```

```

def train_map_estimator(train_images, train_labels):
    """ Inputs: train_images, train_labels
        Returns the MAP estimators theta_map and pi_map """

    # YOU NEED TO WRITE THIS PART
    num_classes = 10
    num_img = train_images.shape[0]
    pixel_per_img = train_images.shape[1]
    train_data = np.where(train_images>0.5,1.0,0.0)

    num_pixel_in_class = np.matmul(train_labels.T,train_data)
    img_in_class = np.sum(train_labels.T,axis=1)
    img_in_class= img_in_class[:,np.newaxis]
    for i in range(0, num_classes):
        num_pixel_in_class[i] = (2+num_pixel_in_class[i])/(4+img_in_class[i][0])
    theta_map = num_pixel_in_class.T

    img_in_class_vec = np.sum(train_labels,axis=0)
    sum_target_all_img = np.sum(np.sum(train_labels,axis=1),axis=0)
    pi_map = img_in_class_vec/sum_target_all_img

    return theta_map, pi_map

def log_likelihood(images, theta, pi):
    """ Inputs: images, theta, pi
        Returns the matrix 'log_like' of loglikelihoods over the input images where
        log_like[i,c] = log p (c |x^(i), theta, pi) using the estimators theta and pi.
        log_like is a matrix of num of images x num of classes
        Note that log likelihood is not only for c^(i), it is for all possible c's. """

    # YOU NEED TO WRITE THIS PART
    num_classes = 10
    sum_img_pixels = np.matmul(images,np.log(theta))
    log_like = np.zeros((num_classes,images.shape[0]))

    for c in range(0,num_classes):
        pi_vec = np.full((1,images.shape[0]),np.log(pi[c]))
        log_like[c] = sum_img_pixels.T[c]+pi_vec

    log_like = log_like.T

    return log_like

def predict(log_like):
    """ Inputs: matrix of log likelihoods
        Returns the predictions based on log likelihood values """

    # YOU NEED TO WRITE THIS PART
    predictions = np.argmax(log_like,axis=1)

```



```
return predictions
```

```
def accuracy(log_like, labels):  
    """ Inputs: matrix of log likelihoods and 1-of-K labels  
    Returns the accuracy based on predictions from log likelihood values"""  
  
    # YOU NEED TO WRITE THIS PART  
    predictions = predict(log_like)  
    num_images = log_like.shape[0]  
    img_idx = np.array(range(num_images))  
    labels_correspondence = labels[img_idx,predictions]  
    matched = np.sum(labels_correspondence)  
    accuracy = matched/num_images  
    return accuracy
```

Q2


a) True : Naïve Assumption : Naïve Bayes assumes that features x_i are conditionally independent given the class c .

b) False: $p(x_i, x_j) = \sum_c p(x_i, x_j | c)$

$$p(x_i)p(x_j) = \sum_c p(x_i | c) \sum_c p(x_j | c)$$

Since $p(x_i, x_j) \neq p(x_i)p(x_j)$,
then they are not independent.

c)




Code for Q2:

```
def image_sampler(theta, pi, num_images):
    """ Inputs: parameters theta and pi, and number of images to sample
    Returns the sampled images"""

    # YOU NEED TO WRITE THIS PART
    theta_index = np.random.choice(len(pi), num_images, p=pi)
    sampled_images = np.random.binomial(1, p=theta.T[theta_index])
    return sampled_images
```


Q3

a)



b) Based on the plot above, I would choose number of eigenvectors as 30 since it has the lowest classification error rate.

c)



Code for Q3:

```

import numpy as np
import matplotlib.pyplot as plt

def load_data(filename, load2=True, load3=True):
    """Loads data for 2's and 3's
    Inputs:
        filename: Name of the file.
        load2: If True, load data for 2's.
        load3: If True, load data for 3's.
    """
    assert (load2 or load3), "Atleast one dataset must be loaded."
    data = np.load(filename)
    if load2 and load3:
        inputs_train = np.hstack((data['train2'], data['train3']))
        inputs_valid = np.hstack((data['valid2'], data['valid3']))
        inputs_test = np.hstack((data['test2'], data['test3']))
        target_train = np.hstack((np.zeros((1, data['train2'].shape[1])), np.ones((1,
data['train3'].shape[1]))))
        target_valid = np.hstack((np.zeros((1, data['valid2'].shape[1])), np.ones((1,
data['valid3'].shape[1]))))
        target_test = np.hstack((np.zeros((1, data['test2'].shape[1])), np.ones((1,
data['test3'].shape[1]))))
    else:
        if load2:
            inputs_train = data['train2']
            target_train = np.zeros((1, data['train2'].shape[1]))
            inputs_valid = data['valid2']
            target_valid = np.zeros((1, data['valid2'].shape[1]))
            inputs_test = data['test2']
            target_test = np.zeros((1, data['test2'].shape[1]))
        else:
            inputs_train = data['train3']
            target_train = np.zeros((1, data['train3'].shape[1]))
            inputs_valid = data['valid3']
            target_valid = np.zeros((1, data['valid3'].shape[1]))
            inputs_test = data['test3']
            target_test = np.zeros((1, data['test3'].shape[1]))

    return inputs_train.T, inputs_valid.T, inputs_test.T, target_train.T, target_valid.T, target_test.T

def l2_distance(a, b):
    """Computes the Euclidean distance matrix between a and b.
    """
    print("a", a.shape)
    print("b", b.shape)
    if a.shape[0] != b.shape[0]:

```

```

    raise ValueError("A and B should be of same dimensionality")

# aa = np.sum(a**2, axis=0)
aa_sqr = np.square(a)
bb_sqr = np.square(b)
aa = np.array(np.sum(aa_sqr, axis=0))[0]
bb = np.array(np.sum(bb_sqr, axis=0))[0]
ab = np.dot(a.T, b)

return np.sqrt(aa[:, np.newaxis] + bb[np.newaxis, :] - 2*ab)

def run_knn(k, train_data, train_labels, valid_data):
    """Uses the supplied training inputs and labels to make
    predictions for validation data using the K-nearest neighbours
    algorithm.

    Note: N_TRAIN is the number of training examples,
          N_VALID is the number of validation examples,
          and M is the number of features per example.

    Inputs:
        k:          The number of neighbours to use for classification
                   of a validation example.
        train_data: The N_TRAIN x M array of training
                   data.
        train_labels: The N_TRAIN x 1 vector of training labels
                     corresponding to the examples in train_data
                     (must be binary).
        valid_data:  The N_VALID x M array of data to
                     predict classes for.

    Outputs:
        valid_labels: The N_VALID x 1 vector of predicted labels
                     for the validation data.
    """
    print("valid_data 2", valid_data.shape)
    print("train_data 2", train_data.shape)

    dist = l2_distance(valid_data.T, train_data.T)
    nearest = np.argsort(dist, axis=1)[:,:k]

    train_labels = train_labels.reshape(-1)
    valid_labels = train_labels[nearest]

    # note this only works for binary labels
    valid_labels = (np.mean(valid_labels, axis=1) >= 0.5).astype(np.int)
    valid_labels = valid_labels.reshape(-1,1)

    return valid_labels

def OneNN(train_inputs, valid_inputs, test_inputs, train_targets, valid_targets, test_targets):
    k_given = [1]

```

```

v_classification_rate = 0
t_classification_rate = 0

for k in k_given:
    print("valid_data1",valid_inputs.shape)
    print("train_data1",train_inputs.shape)
    valid_outputs = run_knn(k,train_inputs,train_targets,valid_inputs)

    if len(valid_outputs) == len(valid_targets):
        count = 0
        for i in range(0, len(valid_outputs)):
            if valid_outputs[i] == valid_targets[i]:
                count += 1
        v_classification_rate=count/len(valid_outputs)

for k in k_given:
    test_outputs = run_knn(k,train_inputs,train_targets,test_inputs)

    if len(test_outputs) == len(test_targets):
        count = 0
        for i in range(0, len(test_outputs)):
            if test_outputs[i] == test_targets[i]:
                count += 1
        t_classification_rate=count/len(test_outputs)
return v_classification_rate,t_classification_rate

def PCA(data, num_principal_comp):
    num_images = data.shape[0]
    num_pixels = data.shape[1]
    mean = np.mean(data,axis=0)
    data_centered = data - np.tile(mean,(num_images,1))
    data_cov = np.cov(data_centered.T)
    eigen_values, eigen_vectors = np.linalg.eig(np.mat(data_cov))
    asc_sorted_eigval_index = np.argsort(eigen_values)
    desc_sorted_eigval_index = asc_sorted_eigval_index[-(num_principal_comp+1):-1]
    eigvec_chosen = eigen_vectors[:,desc_sorted_eigval_index]

    low_d_data = np.matmul(data_centered,eigvec_chosen)
    reconstructed_data = np.matmul(low_d_data, eigvec_chosen.T) + mean
    return reconstructed_data, low_d_data

if __name__ == '__main__':
    inputs_train, inputs_valid, inputs_test, target_train, target_valid, target_test = load_data('./
digits.npz')
    pca_lst = [2,5,10,20,30]
    valid_rates = []
    test_rates = []
    for pca in pca_lst:
        reconstructed_data_train, low_d_data_train = PCA(inputs_train, pca)
        reconstructed_data_valid, low_d_data_valid = PCA(inputs_valid, pca)
        reconstructed_data_test, low_d_data_test = PCA(inputs_test, pca)
        v_rate, t_rate = OneNN(low_d_data_train,
                                low_d_data_valid,
                                low_d_data_test,

```

```
        target_train,
        target_valid,
        target_test)
    valid_rates.append(1-v_rate)
    test_rates.append(1-t_rate)
plt.plot(pca_lst,valid_rates,label='valid_set')
plt.xlabel("K")
plt.ylabel("v_classification_error_rate")
plt.title("Validation error rate VS K")
plt.savefig("PCA validation error rate.png")

plt.clf()
plt.plot(pca_lst,test_rates,label='test_set')
plt.xlabel("K")
plt.ylabel("t_classification_error_rate")
plt.title("Test error rate VS K")
plt.savefig("PCA test error rate.png")
plt.clf()
```